

DATA

The data was collected from the finance library “yfinance” in python.

We are studying data from 1981 to 2020 and will forecast for 2021.

```
# using Yahoo finance API to get stock data.  
stocksdata="AAPL"  
df=yf.download(stocksdata,start="1981-1-1",end="2020-12-31")
```

df


	Date	Open	High	Low	Close	Adj Close	Volume
0	1981-01-02	0.154018	0.155134	0.154018	0.154018	0.119849	21660800
1	1981-01-05	0.151228	0.151228	0.150670	0.150670	0.117244	35728000
2	1981-01-06	0.144531	0.144531	0.143973	0.143973	0.112032	45158400
3	1981-01-07	0.138393	0.138393	0.137835	0.137835	0.107256	55686400
4	1981-01-08	0.135603	0.135603	0.135045	0.135045	0.105085	39827200
...
10081	2020-12-23	132.160004	132.429993	130.779999	130.960007	129.406555	88223700
10082	2020-12-24	131.320007	133.460007	131.100006	131.970001	130.404572	54930100
10083	2020-12-28	133.990005	137.339996	133.509995	136.690002	135.068588	124486200
10084	2020-12-29	138.050003	138.789993	134.339996	134.869995	133.270187	121047300
10085	2020-12-30	135.580002	135.990005	133.399994	133.720001	132.133820	96452100

10086 rows × 7 columns

Predictions using LR:

[Lr of sf adaniport.ipynb - Colaboratory \(google.com\)](#)

R2-score: 0.9913189150766206



Predictions Using LSTM:

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 100, 50)	10400
lstm_1 (LSTM)	(None, 100, 50)	20200
lstm_2 (LSTM)	(None, 50)	20200
dense (Dense)	(None, 1)	51

=====
Total params: 50,851

Trainable params: 50,851

Non-trainable params: 0

LSTM ARCHITECTURE:


```
model=Sequential()
model.add(LSTM(50,return_sequences=True,input_shape=(X_train.shape[1],1)))
model.add(LSTM(50,return_sequences=True))
model.add(LSTM(50))
model.add(Dense(1,activation='linear'))
model.compile(loss='mean_squared_error',optimizer='adam')
```

APPLE


R2 score : 0.9820609540002779

MSE : 0.06042422772480179

Stock forecasting apple.ipynb - Colaboratory (google.com)




Forecasting USING LSTM:



GOLDMAN SACHS

R2-Score - 0.9762633409794264

[Stock forecasting GS.ipynb - Colaboratory \(google.com\)](#)



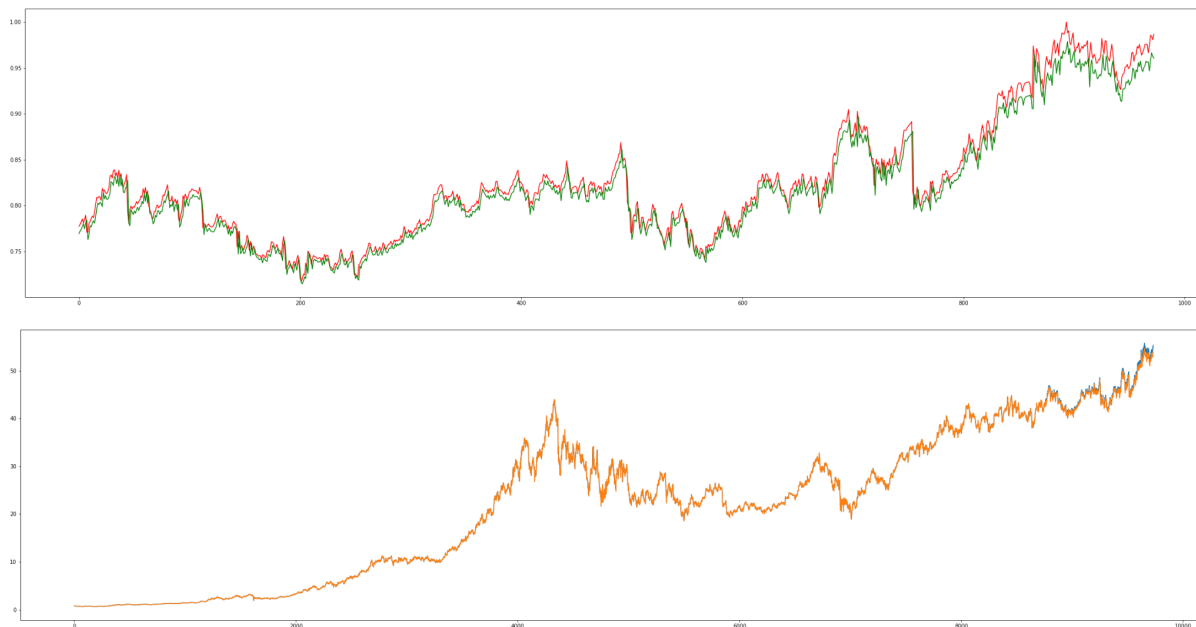
Forecasting USING LSTM:




COCA-COLA

R2-Score-0.9630954002841426

[Stock forecasting coca-cola.ipynb - Colaboratory \(google.com\)](#)




Forecasting USING LSTM:



IBM

R2-Score-0.978463002841426

[Stock forecasting IBM.ipynb - Colaboratory \(google.com\)](#)



USING TRANSFORMERS

[APPLE_STOCKS\(moving_average\).ipynb](#)

Time Embeddings:

Transformers combined with natural language data tend to utilize positional encoding to provide the model with word order. In detail, the positional encoding represents a word's value and position in a sentence, allowing the Transformer to obtain knowledge about a sentence structure and word inter-dependencies, so for the information of time to go through the network, we will use TIME2VEC.

Time to vector:

```
class Time2Vector(Layer):
    def __init__(self, seq_len, **kwargs):
        super(Time2Vector, self).__init__()
        self.seq_len = seq_len

    def build(self, input_shape):
        '''Initialize weights and biases with shape (batch, seq_len)'''
        self.weights_linear = self.add_weight(name='weight_linear',
                                              shape=(int(self.seq_len),),
                                              initializer='uniform',
                                              trainable=True)

        self.bias_linear = self.add_weight(name='bias_linear',
                                           shape=(int(self.seq_len),),
                                           initializer='uniform',
                                           trainable=True)

        self.weights_periodic = self.add_weight(name='weight_periodic',
                                                shape=(int(self.seq_len),),
                                                initializer='uniform',
                                                trainable=True)

        self.bias_periodic = self.add_weight(name='bias_periodic',
                                             shape=(int(self.seq_len),),
                                             initializer='uniform',
                                             trainable=True)

    def call(self, x):
        '''Calculate linear and periodic time features'''
        x = tf.math.reduce_mean(x[:, :, :4], axis=-1)
        time_linear = self.weights_linear * x + self.bias_linear # Linear time feature
        time_linear = tf.expand_dims(time_linear, axis=-1) # Add dimension (batch, seq_len, 1)

        time_periodic = tf.math.sin(tf.multiply(x, self.weights_periodic) + self.bias_periodic)
        time_periodic = tf.expand_dims(time_periodic, axis=-1) # Add dimension (batch, seq_len, 1)
        return tf.concat([time_linear, time_periodic], axis=-1) # shape = (batch, seq_len, 2)

    def get_config(self): # Needed for saving and loading model with custom Layer
        config = super().get_config().copy()
        config.update({'seq_len': self.seq_len})
        return config
```



Transformer Architecture:

Model: "model_5"

Layer (type)	Output Shape	Param #	Connected to
input_6 (InputLayer)	[(None, 128, 5)]	0	[]
time2_vector_5 (Time2Vector)	(None, 128, 2)	512	['input_6[0][0]']
concatenate_5 (Concatenate)	(None, 128, 7)	0	['input_6[0][0]', 'time2_vector_5[0][0]']
transformer_encoder_15 (TransformerEncoder)	(None, 128, 7)	99114	['concatenate_5[0][0]', 'concatenate_5[0][0]', 'concatenate_5[0][0]']
transformer_encoder_16 (TransformerEncoder)	(None, 128, 7)	99114	['transformer_encoder_15[0][0]', 'transformer_encoder_15[0][0]', 'transformer_encoder_15[0][0]']
transformer_encoder_17 (TransformerEncoder)	(None, 128, 7)	99114	['transformer_encoder_16[0][0]', 'transformer_encoder_16[0][0]', 'transformer_encoder_16[0][0]']
global_average_pooling1d_5 (GlobalAveragePooling1D)	(None, 128)	0	['transformer_encoder_17[0][0]']
dropout_10 (Dropout)	(None, 128)	0	['global_average_pooling1d_5[0][0]']
dense_10 (Dense)	(None, 64)	8256	['dropout_10[0][0]']
dropout_11 (Dropout)	(None, 64)	0	['dense_10[0][0]']
dense_11 (Dense)	(None, 1)	65	['dropout_11[0][0]']
=====			
Total params: 306,175			
Trainable params: 306,175			
Non-trainable params: 0			

RESULT OF TRANSFORMER MODEL:



Transformer + TimeEmbedding Model


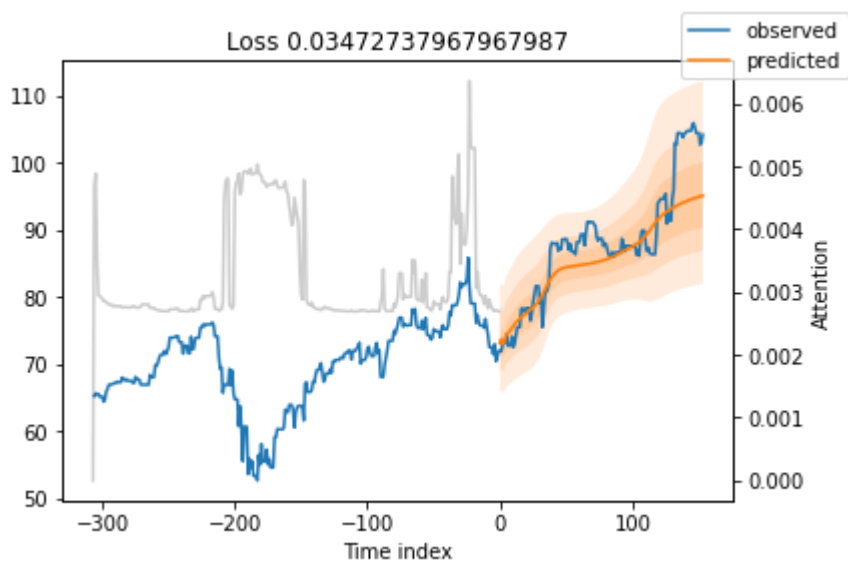



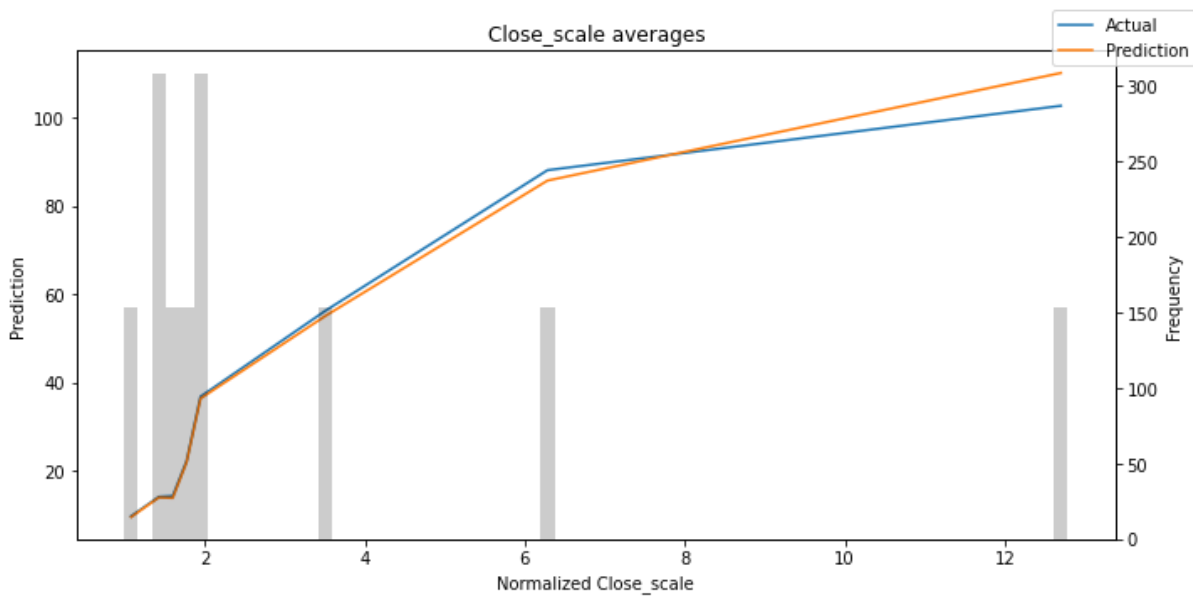


Ig, we can conclude that this transformer model is shit.

TEMPORAL FUSION TRANSFORMER

We trained the model to forecast the values of stock closing price for 154 days based on the given context data of 308 days.









PROPHET - Facebook's Time Series Forecasting model

[forecasting using prophet.ipynb - Colaboratory \(google.com\)](#)

We trained Prophet on apples stocks from 1998 to 2020.



We got following results in forecasting for the year 2021.



We used the Base model as it was not showing much improvement on hyperparameter tuning.

```
[ ] mape(pred1['yhat'],df_real['Close'])
```

```
0.44282468919904
```

```
▶ min(tuning_results['mape'])
```

```
↳ 0.43072457760166066
```

Forecasting for One Month.

